

# Shopping Lists on the Cloud

A Scalable, Highly Available Distributed System

Emanuel Gestosa, José Gaspar, Lucas Sousa

SDLE, MEIC - FEUP

December 9, 2023

# Outline

## 1 Introduction

## 2 Architecture

## 3 Implementation

- Communication
- Node preference lists
- Replication
- Conflict-free Replicated Data Types (CRDTs)
- Local
- Fault tolerance

## 4 Discussion

## 5 Conclusion

# Introduction

# Introduction

We want to build a system that can manage shopping lists from different users. Our system should follow the following guidelines:

- ▶ Local-first application.
  - Store data locally first and on the cloud.
- ▶ Share lists between clients.
- ▶ Version conflicts solved with CRDTs.
- ▶ High availability.

# Architecture

# Architecture

- ▶ **Local-first application:** Provides an interface between the client and the cloud, but can also be accessed while offline.
- ▶ **Proxy:** Gateway between the local application and the rest of the cloud infrastructure.
- ▶ **Node:** Deals with requests to read or update shopping lists, as well as data replication.
- ▶ **Coordinator:** A special kind of Node, that receives requests directly from the client, and may use other Nodes to store/retrieve replicas.

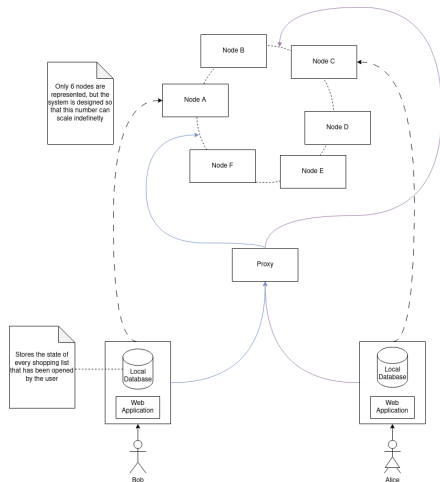


Figure: High-level overview of the design  
Architecture

# Implementation

# Implementation

## Communication

- ▶ Communication is done via HTTP and with ZeroMQ.
- ▶ We have communication between the following entities:
  - Proxy - Local application: to inform the local application about the Coordinator of the desired list.
  - Proxy - Node: to inform the Node about its preference list and for a Node to join/leave the circle.
  - Node - Node: to manage replicas.
  - Node - Local application: to retrieve/update lists from the cloud.



# Implementation

## Node preference lists

- ▶ A Node joins the circle by sending an HTTP POST request to the Proxy.
- ▶ The Proxy updates the Node circle, and the preference lists of the Nodes.
- ▶ The Proxy publishes the updated preference lists.
- ▶ Every Node subscribes to updates to their preference lists, and checks for updates every time they need to act as Coordinator.
- ▶ A Coordinator will use its preference list to know which Nodes it should use for consulting or updating replicas.

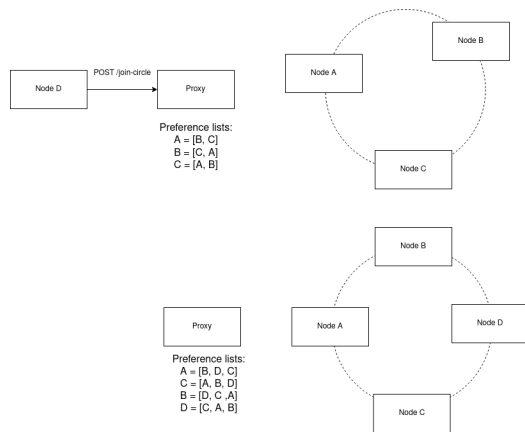


Figure: New Node joining circle

# Implementation

## Replication

### Quorum strategy

- ▶ Coordinator has it's own preference list.
- ▶ The Coordinator waits for a request from the client.
- ▶ if PUT:
  - Coordinator stores information in its database.
  - Sends request for all nodes in its preference list for replica creation.
  - Waits for responses.
  - PUT always succeeds in ensuring high availability.
- ▶ if GET:
  - Coordinator sends requests for all nodes in its preference list for getting existent replicas.
  - Waits for responses.
  - GET only succeeds if at least one replica is received.
  - Merge replicas.
  - Return merged replicas to the Client.

# Implementation

## Resolution of version conflicts - Conflict-free Replicated Data Types (CRDTs)

Operation-Based CRDTs were used:

- ▶ All operations are saved as an operation entry, with an id hash, item name, type field ("ADD", "REMOVE"), and a count field.
- ▶ Both the add and remove operations are treated the same way, the only difference is how they are displayed in the GUI.
- ▶ Store operations over the list.
  - Add N element to a list.
  - Remove N element from a list.
- ▶ Merge versions function:
  - Used GSet strategy.
    - ▶ Node merging by adding all operations to a set, thereby keeping one entry of every operation.
- ▶ Inspired by git commits.

# Implementation

## Local

- ▶ Django web app, assumes that the Django Backend and Frontend are running on the same machine, emulating a local application.
- ▶ On page load, queries the proxy for the IP address and port of the appropriate node.
- ▶ Subsequently queries the node, sending its locally saved items and retrieving the merge of all items.
- ▶ Sends a PUT request when making changes and a GET request when retrieving from the server.
- ▶ At 5-second intervals polls the nodes for changes.
- ▶ Saves all changes locally, regardless of node connection.

# Implementation

## Local

### Available Lists

Proxy IP: localhost	<input type="text"/>	Set IP
Hash:	<input type="text"/>	Add Existing List
Title:	<input type="text"/>	Create New List

<a href="#">MyList1</a>	<input type="button" value="Remove"/>	<input type="button" value="Share"/>
-------------------------	---------------------------------------	--------------------------------------

Figure: Main Page with Lists

### List: MyList1

☒ Online

Item Title:	<input type="text"/>	<input type="button" value="Add New Item"/>
-------------	----------------------	---

Banana	<input type="text" value="3"/>
Orange	<input type="text" value="1"/>
Apple	<input type="text" value="1"/>
Pear	<input type="text" value="1"/>
Tangerine	<input type="text" value="1"/>

Figure: List Page with Items

# Implementation

What if something goes wrong?

- ▶ What if a node is not responsive?
  - The Coordinator knows it upon GET or PUT requests.
  - Informs proxy about it and proxy disconnects it from the circle.
  - The system works fine because of the replicas on other nodes from the coordinator's preference list.
- ▶ Failure points:
  - Proxy failure.
  - All nodes from a preference list are not responsive.

# Discussion

# Discussion

## CAP Theorem

- ▶ On our project, we wanted to have a highly available system.
- ▶ That is why we allow the PUT requests to always succeed, and GET requests succeed with just one replica.
- ▶ To achieve this we cannot guarantee that replicas are consistent all the time, but on the other hand, if a node goes down, its information should still be saved somewhere else, and thus always available.
- ▶ However, the synchronization of replicas is done upon GET requests.



## Conclusion

# Conclusion

- ▶ We fulfilled our goal of achieving a highly available distributed system.
- ▶ Users can create and share lists, and work on them at the same time.
- ▶ Conflicts are always resolved, and data loss is highly unlikely.
- ▶ Possible improvements:
  - Addition of backup proxies to avoid the system being down on the chance the Proxy fails.
  - Local application currently uses polling, an improvement would be to use a pub/sub approach.